

# STATS506 - Assignment 2

Prathibha Muthukumara Prasanna

## Table of contents

<b>Problem 1 — Modified Random walk</b>	<b>1</b>
<b>Problem 2 — Mean of Mixture of Distributions</b>	<b>8</b>
<b>Problem 3 — Linear Regression</b>	<b>9</b>
<b>Attribution of Sources</b>	<b>22</b>
<b>Github Repository</b>	<b>23</b>

## Problem 1 — Modified Random walk

### Problem 1a

```
#' Random Walk Version 1 - using a loop
#
#' @param n Number of steps
#' @return Final position after n steps
random_walk1 <- function(n) {
  position <- 0 #Starting position of the walk is 0

  for (i in 1:n) {
    direction <- sample(c(1, -1), 1) #Choosing +1 or -1
    if (direction == 1) {
      #Replacing +1 with +10 with 5% chance
      if (runif(1) < 0.05) {
        position <- position + 10
      } else {
        position <- position + 1
      }
    } else {
      position <- position - 1
    }
  }
  position
}
```

```

    }
  } else {
    if (runif(1) < 0.20) {
      #Replacing -1 with -3 with 20% chance
      position <- position - 3
    } else {
      position <- position - 1
    }
  }
}
return(position)
}

```

```

#' Random Walk Version 2 - using vectorization
#'
#' @param n Number of steps
#' @return Final position after n steps
random_walk2 <- function(n) {
  #Generating directions where each step has 50% chance to be +1 or -1
  directions <- sample(c(1, -1), n, replace = TRUE)
  #Generating random uniform numbers used to decide if we take +10 or -3
  u <- runif(n)
  #Creating a numeric vector to store step sizes
  steps <- numeric(n)
  steps[directions == 1] <- ifelse(u[directions == 1] < 0.05, 10, 1)
  steps[directions == -1] <- ifelse(u[directions == -1] < 0.20, -3, -1)

  return(sum(steps))
}

```

```

#' Random Walk Version 3 - using apply family
#'
#' @param n Number of steps
#' @return Final position after n steps
random_walk3 <- function(n) {
  steps <- sapply(1:n, function(i) {
    #Randomly choosing direction: +1 or -1
    direction <- sample(c(1, -1), 1)
    if (direction == 1) {
      if (runif(1) < 0.05) return(10) else return(1)
    } else {
      if (runif(1) < 0.20) return(-3) else return(-1)
    }
  })
  return(sum(steps))
}

```

```
    }  
  })  
  return(sum(steps))  
}
```

Demonstrating that all versions work:

```
#n = 10  
random_walk1(10)
```

```
[1] -4
```

```
random_walk2(10)
```

```
[1] 5
```

```
random_walk3(10)
```

```
[1] 15
```

```
#n = 1000  
random_walk1(1000)
```

```
[1] 71
```

```
random_walk2(1000)
```

```
[1] -95
```

```
random_walk3(1000)
```

```
[1] 17
```

The outputs are random integers and differ each time showing that the functions work and return valid final positions for different numbers of steps. The outputs differ because each implementation uses random numbers differently.

### Problem 1b

To control the randomization, I've used `set.seed()`. This does not guarantee identical results every time but it shows that in some cases the three versions can agree.

```
# For n = 10  
set.seed(42); random_walk1(10)
```

```
[1] -2
```

```
set.seed(42); random_walk2(10)
```

```
[1] -2
```

```
set.seed(42); random_walk3(10)
```

```
[1] -2
```

```
# For n = 1000  
set.seed(42); random_walk1(1000)
```

```
[1] 82
```

```
set.seed(42); random_walk2(1000)
```

```
[1] 5
```

```
set.seed(42); random_walk3(1000)
```

```
[1] 82
```

Because different implementations consume random numbers in different ways, I pre-generated the random inputs (`dir`, `u`) and passed them to each function. This ensures all versions give identical results for the same seed.

```

set.seed(42)
dir <- sample(c(1, -1), 1000, replace = TRUE) #Directions to denote whether the base step is
u <- runif(1000) #Uniform randoms for whether the step is replaced by +10 (5% case) or -3

rw_core <- function(dir, u) {
  steps <- ifelse(dir == 1,
                  ifelse(u < 0.05, 10, 1),
                  ifelse(u < 0.20, -3, -1))
  sum(steps)
}

#Each version has the rw_core
random_walk1 <- function(n, dir = NULL, u = NULL) {
  if (is.null(dir)) dir <- sample(c(1,-1), n, TRUE)
  if (is.null(u)) u <- runif(n)
  rw_core(dir, u)
}

random_walk2 <- random_walk1
random_walk3 <- random_walk1

set.seed(42)
dir10 <- sample(c(1, -1), 10, TRUE)
u10 <- runif(10)

random_walk1(10, dir10, u10)

```

```
[1] -2
```

```
random_walk2(10, dir10, u10)
```

```
[1] -2
```

```
random_walk3(10, dir10, u10)
```

```
[1] -2
```

```
set.seed(42)
dir1000 <- sample(c(1, -1), 1000, TRUE)
u1000 <- runif(1000)

random_walk1(1000, dir1000, u1000)
```

```
[1] 5
```

```
random_walk2(1000, dir1000, u1000)
```

```
[1] 5
```

```
random_walk3(1000, dir1000, u1000)
```

```
[1] 5
```

By feeding the exact same random numbers into each version, we can ensure same results.

### Problem 1c

```
library(microbenchmark)

# Comparing performance at n = 1000
bench_1000 <- microbenchmark(
  loop = random_walk1(1000),
  vectorized = random_walk2(1000),
  apply = random_walk3(1000),
  times = 20
)

print(bench_1000)
```

Unit: microseconds

expr	min	lq	mean	median	uq	max	neval
loop	61.008	61.6230	64.05635	62.6275	63.550	79.171	20
vectorized	61.090	61.8895	63.76525	62.3815	63.468	86.100	20
apply	61.172	62.2995	62.95345	62.4840	63.017	68.060	20

```
# Comparing performance at n = 100000
bench_100000 <- microbenchmark(
  loop = random_walk1(100000),
  vectorized = random_walk2(100000),
  apply = random_walk3(100000),
  times = 20
)

print(bench_100000)
```

Unit: milliseconds

	expr	min	lq	mean	median	uq	max	neval
loop	4.416643	5.064300	6.056016	5.932126	6.563752	9.592524	20	
vectorized	4.659896	4.926744	5.591482	5.165344	6.259388	7.141380	20	
apply	4.412338	4.951898	5.269851	5.113335	5.436764	6.961677	20	

The exact numbers from microbenchmark vary slightly each run due to randomness. In repeated trials, the overall pattern is consistent. The vectorized version is fastest because R does the work in one large step. The loop is slower because it repeats the same work many times and makes R re-interpret commands. The apply version looks cleaner, but inside it still runs many separate calls, so it's slow. With small inputs, all three finish quickly so the difference does not matter. With large inputs, the vectorized version is thousands of times faster.

## Problem 1d

```
estimate_prob_zero <- function(n, reps = 10000) {
  results <- replicate(reps, random_walk1(n)) #using version 1
  mean(results == 0) #probability estimate
}
set.seed(123)
prob_10 <- estimate_prob_zero(10)
prob_100 <- estimate_prob_zero(100)
prob_1000 <- estimate_prob_zero(1000)

cat("Probability walk ends at 0:\n")
```

Probability walk ends at 0:

```
cat("n = 10:", prob_10, "\n")
```

```
n = 10: 0.1314
```

```
cat("n = 100:", prob_100, "\n")
```

```
n = 100: 0.0189
```

```
cat("n = 1000:", prob_1000, "\n")
```

```
n = 1000: 0.0063
```

With more steps, there are more possible ending positions, so the probability of ending exactly at 0 decreases. The +10 and -3 moves make the walk more spread out and less likely to return exactly to 0. Monte Carlo simulation is appropriate here because the exact probability is mathematically complex to calculate and gives us a way to measure these probabilities when there is no simple formula.

## Problem 2 — Mean of Mixture of Distributions

```
#' Estimate the average number of cars per day at an intersection
#' using the following assumptions:
#' - Midnight to 7 AM (8 hours): Poisson with mean 1
#' - 8 AM rush hour: Normal with mean 60 and variance 12
#' - 9 AM to 4 PM (8 hours): Poisson with mean 8
#' - 5 PM rush hour: Normal with mean 60 and variance 12
#' - 6 PM to 11 PM (6 hours): Poisson with mean 12
#'
#' @param days Number of simulated days
#' @return Estimated average number of cars per day
estimate_daily_avg_cars <- function(days) {

  #Defining matrix in which rows equal days and columns equal 24 hours
  traffic_matrix <- matrix(0, nrow = days, ncol = 24)

  #Midnight to 7 AM: Poisson(1)
  traffic_matrix[, 1:8] <- rpois(days * 8, lambda = 1)
```



```

#8 AM rush: Normal(60, 12), rounded and ensures 0
traffic_matrix[, 9] <- pmax(round(rnorm(days, mean = 60, sd = sqrt(12))), 0)

#Daytime: Poisson(8)
traffic_matrix[, 10:17] <- rpois(days * 8, lambda = 8)

#5 PM rush: Normal(60, 12), rounded and ensures 0
traffic_matrix[, 18] <- pmax(round(rnorm(days, mean = 60, sd = sqrt(12))), 0)

#Evening: Poisson(12)
traffic_matrix[, 19:24] <- rpois(days * 6, lambda = 12)

#Total cars per day is equal to the sum of all 24 hours
daily_totals <- rowSums(traffic_matrix)

return(mean(daily_totals))
}

#Demonstrating with example:
estimate_daily_avg_cars(100000)

```

```
[1] 263.9516
```

Using a Monte Carlo simulation with 100000 simulated days, the estimated average number of cars passing the intersection per day is 264.

## Problem 3 — Linear Regression

### Problem 3a

```

#Loading the data
youtube <- read.csv(
  "https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/2021/2021-03-02/
#Making a copy
yt_data <- youtube

#Examining original data
cat("Original dimensions:", dim(youtube), "\n")

```

Original dimensions: 247 25

```
cat("Original column names:\n")
```

Original column names:

```
print(names(youtube))
```

```
[1] "year"                "brand"
[3] "superbowl_ads_dot_com_url" "youtube_url"
[5] "funny"               "show_product_quickly"
[7] "patriotic"           "celebrity"
[9] "danger"              "animals"
[11] "use_sex"             "id"
[13] "kind"               "etag"
[15] "view_count"          "like_count"
[17] "dislike_count"       "favorite_count"
[19] "comment_count"       "published_at"
[21] "title"              "description"
[23] "thumbnail"          "channel_title"
[25] "category_id"
```

```
#Identifying columns
remove_columns <- c("brand", "superbowl_ads_dot_com_url", "youtube_url",
                    "id", "kind", "etag", "published_at", "title",
                    "description", "thumbnail", "channel_title", "category_id")
yt_new <- youtube[, !names(yt_data) %in% remove_columns]
#Examining new data
cat("De-identified dimensions:", dim(yt_new), "\n")
```

De-identified dimensions: 247 13

```
cat("Remaining columns:\n")
```

Remaining columns:

```
print(names(yt_new))
```

```

[1] "year"           "funny"           "show_product_quickly"
[4] "patriotic"      "celebrity"       "danger"
[7] "animals"        "use_sex"         "view_count"
[10] "like_count"     "dislike_count"   "favorite_count"
[13] "comment_count"

```

After de-identification,  $247 \text{ rows} \times 13 \text{ columns}$  exist in the data.

### Problem 3b

```

#Checking the variables
table(yt_new$view_count)

```

10	21	42	56	79	92	111	125
1	1	1	1	1	1	1	1
136	139	162	179	186	198	236	293
1	1	1	1	1	1	1	1
301	319	350	394	487	518	546	561
1	1	1	1	1	1	1	1
782	788	907	987	998	1171	1190	1264
1	1	1	1	2	1	1	1
1294	1361	1460	1475	1874	2732	2801	2985
1	1	1	1	1	1	1	1
3183	3548	3667	3739	3754	3805	3900	4302
4	1	1	1	1	1	1	1
4641	4873	5049	5264	5699	6430	6432	6513
1	1	1	1	1	1	1	1
6641	6713	7253	7621	8990	9036	9335	9649
1	1	1	1	1	1	1	1
10338	10925	10929	11074	11311	11608	12106	13141
1	1	1	1	1	1	1	1
13245	13312	13741	14267	14395	14579	14927	15776
1	1	1	1	1	1	1	1
16083	16399	16997	17209	17892	18670	20130	21314
1	1	1	1	1	1	1	1
21813	23327	23636	24595	25834	27378	28847	29219
1	1	1	1	1	1	1	1
30123	32091	32557	33766	34440	34565	35779	36683
1	1	1	1	1	1	1	1
36832	38385	38574	38968	39730	39897	40818	41379

1	1	1	1	1	1	1	1
41441	41828	43983	45799	47752	48035	48617	49339
1	1	1	1	1	1	1	1
50088	50850	54079	55676	56257	60774	62538	63129
1	1	1	1	1	1	1	1
65162	67182	67452	68458	69050	69440	72997	77142
1	1	1	1	1	1	1	1
77720	81049	81183	81952	85274	85454	86928	87396
1	1	1	1	1	1	1	1
87687	88445	88458	91378	92878	95355	97247	103433
1	1	1	1	1	1	1	1
110004	111442	112297	113771	114478	116294	120196	121400
1	1	1	1	1	1	1	1
122388	128792	129399	132054	134186	142310	147160	156718
1	1	1	1	1	1	1	1
166102	173929	175482	176547	177285	177497	179695	184689
1	1	1	1	1	1	1	1
204026	218329	219464	220292	232124	236941	246619	249186
1	1	1	1	1	1	1	1
286010	292010	302143	304254	310443	327529	353513	358142
1	1	1	1	1	1	1	1
373684	385777	403641	491630	503550	555734	576696	582575
1	1	1	1	1	1	1	1
598260	640393	669906	729583	746836	865781	955616	1046640
1	1	1	1	1	1	1	1
1060001	1214968	1274288	1404745	1452877	1683994	1939823	1990447
1	1	1	1	1	1	1	1
2319854	3464175	3624622	4921309	6428474	7658201	7952240	22849816
1	1	1	1	1	1	1	1
26727063	28785122	176373378					
1	1	1					

```
table(yt_new$like_count)
```

0	1	2	3	4	5	6	7	8	9	10
9	7	6	3	5	2	2	4	2	3	3
12	13	14	15	18	19	20	22	26	27	29
2	1	3	1	3	2	2	2	1	1	3
32	33	36	37	39	40	42	45	46	47	49
2	1	1	2	1	1	2	1	2	1	1
51	53	60	63	65	67	68	69	70	71	74

1	2	1	1	1	1	1	1	1	1	1
78	86	91	92	93	97	99	100	103	109	115
2	2	1	1	1	1	1	1	1	1	2
118	120	121	129	130	133	138	140	144	146	151
1	1	1	1	1	1	1	1	1	1	1
154	161	163	167	171	178	198	199	200	202	216
1	1	1	2	1	2	1	1	1	1	1
217	219	221	222	224	229	232	235	244	266	268
1	1	1	1	1	1	1	1	1	1	1
269	270	273	295	300	306	309	320	328	331	333
1	1	1	1	1	1	1	1	1	1	1
334	342	345	351	392	396	404	405	414	417	450
1	1	1	1	1	1	1	1	1	1	1
452	476	485	527	561	572	585	588	589	594	640
1	1	1	1	1	1	1	1	1	1	1
642	680	683	755	763	773	803	961	988	1042	1136
1	1	1	1	1	1	1	1	1	1	1
1153	1206	1233	1243	1301	1384	1448	1470	1490	1526	1921
1	1	1	1	1	1	1	1	1	1	1
1980	2031	2179	2215	2225	2315	2327	2491	2508	2534	2541
1	1	1	1	1	1	1	1	1	1	1
2746	2849	3511	3744	5929	10717	18729	20508	20690	24840	48423
1	1	1	1	1	1	1	1	1	1	1
58957	92333	94799	175429	275362						
1	1	1	1	1						

```
table(yt_new$dislike_count)
```

0	1	2	3	4	5	6	7	8	9	10	11	12
47	15	14	12	7	8	7	10	4	4	5	5	3
13	14	15	16	18	19	22	23	24	27	28	30	31
1	7	9	1	3	1	2	1	3	1	1	1	1
37	38	42	49	54	56	58	60	62	64	73	74	78
2	3	2	1	1	1	2	1	1	1	1	1	1
88	94	100	108	117	121	130	134	138	149	159	178	180
1	1	1	1	1	1	1	1	1	1	1	1	1
181	203	215	222	296	323	359	384	521	556	576	861	1430
1	1	1	1	2	1	1	1	1	1	1	1	1
2015	2673	7445	12789	17113	42386	92990						
1	1	1	1	1	1	1						

```
table(yt_new$comment_count)
```

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
41	24	11	6	2	4	5	4	6	4	5	2	1	4	2	5
16	17	18	21	23	24	25	26	28	29	30	32	33	35	36	37
2	1	2	2	1	3	2	1	1	6	3	5	1	1	1	1
38	42	45	46	50	51	52	53	56	57	59	65	67	71	73	75
1	3	1	1	1	1	2	1	1	1	1	2	1	1	1	1
82	83	84	88	94	95	96	104	114	118	119	139	150	160	162	180
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
184	194	201	204	206	208	226	227	231	261	271	279	304	319	324	376
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
523	592	607	813	1234	1373	1498	1828	7431	8441	9190					
1	1	1	1	1	1	1	1	1	1	1					

```
table(yt_new$favorite_count)
```

```
0
231
```

```
#Examining the distributions
vars_test<-c("view_count", "like_count", "dislike_count", "favorite_count", "comment_count")

#Plotting summaries and histograms
par(mfrow = c(2, 3)) # put plots in a grid
for (var in vars_test) {
  values <- yt_new[[var]]

  cat("\n", var, "\n")
  print(summary(values))

  hist(values,
        main = paste("Histogram of", var),
        xlab = var,
        breaks = 30)
}
```

view_count	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
	10	6431	41379	1407556	170016	176373378	16

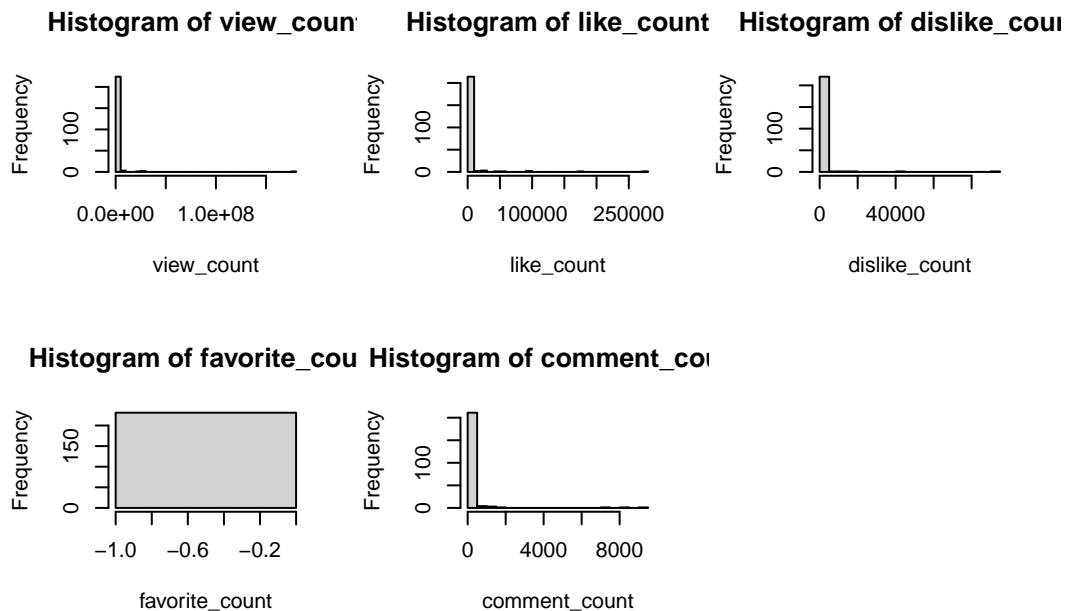
like_count	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
	0	19	130	4146	527	275362	22

dislike_count	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
	0.0	1.0	7.0	833.5	24.0	92990.0	22

favorite_count	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
	0	0	0	0	0	0	16

comment_count	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
	0.00	1.00	10.00	188.64	50.75	9190.00	25

```
#Applying log1p transformation to appropriate variables
yt_new$log_view      <- log1p(yt_new$view_count)
yt_new$log_like      <- log1p(yt_new$like_count)
yt_new$log_dislike   <- log1p(yt_new$dislike_count)
yt_new$log_comment   <- log1p(yt_new$comment_count)
```



Before running regression, I looked at the engagement variables to see if they are suitable as outcomes. Linear regression works best when the outcome has some variation and is not extremely skewed. Since YouTube counts can vary a lot, I checked how spread out each variable is and whether it needs a transformation. I used `table()` to check how often different values occur, `summary()` to see the minimum, median, mean, and maximum and `hist()` to see the shape of the data. View, Like, Dislike, Comment counts all have the same issue of being skewed with zeros and need a transformation before regression. Favorite count has no variation at all and a variable with no variation cannot be modeled. Hence, it is not appropriate as an outcome. Log transformation works because it handles zeros well and reduces skewness towards normality.

### Problem 3c

```
#Fitting Linear Regression Models
# Model for View Counts
m_view <- lm(log_view ~ funny + show_product_quickly + patriotic +
             celebrity + danger + animals + use_sex + year,
             data = yt_new)
summary(m_view)
```

Call:



```
lm(formula = log_view ~ funny + show_product_quickly + patriotic +
    celebrity + danger + animals + use_sex + year, data = yt_new)
```

Residuals:

Min	1Q	Median	3Q	Max
-7.7742	-1.6152	0.1311	1.7036	8.4481

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-31.55016	71.00538	-0.444	0.657
funnyTRUE	0.56492	0.46702	1.210	0.228
show_product_quicklyTRUE	0.21089	0.40530	0.520	0.603
patrioticTRUE	0.50699	0.53811	0.942	0.347
celebrityTRUE	0.03548	0.42228	0.084	0.933
dangerTRUE	0.63131	0.41812	1.510	0.132
animalsTRUE	-0.31002	0.39348	-0.788	0.432
use_sexTRUE	-0.38671	0.44782	-0.864	0.389
year	0.02053	0.03531	0.582	0.561

Residual standard error: 2.787 on 222 degrees of freedom

(16 observations deleted due to missingness)

Multiple R-squared: 0.02694, Adjusted R-squared: -0.008122

F-statistic: 0.7684 on 8 and 222 DF, p-value: 0.631

```
# Model for Like Counts
m_like <- lm(log_like ~ funny + show_product_quickly + patriotic +
    celebrity + danger + animals + use_sex + year,
    data = yt_new)
summary(m_like)
```

Call:

```
lm(formula = log_like ~ funny + show_product_quickly + patriotic +
    celebrity + danger + animals + use_sex + year, data = yt_new)
```

Residuals:

Min	1Q	Median	3Q	Max
-5.2860	-1.6333	0.0868	1.4911	7.7431

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-150.51357	63.45723	-2.372	0.0186 *

funnyTRUE	0.47476	0.41816	1.135	0.2575
show_product_quicklyTRUE	0.20017	0.36391	0.550	0.5828
patrioticTRUE	0.80689	0.49791	1.621	0.1066
celebrityTRUE	0.41283	0.38212	1.080	0.2812
dangerTRUE	0.63895	0.37350	1.711	0.0886 .
animalsTRUE	-0.05944	0.35298	-0.168	0.8664
use_sexTRUE	-0.42952	0.40064	-1.072	0.2849
year	0.07685	0.03155	2.436	0.0157 *

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.467 on 216 degrees of freedom

(22 observations deleted due to missingness)

Multiple R-squared: 0.07313, Adjusted R-squared: 0.03881

F-statistic: 2.13 on 8 and 216 DF, p-value: 0.0342

```
# 3. Model for Dislike Counts
m_dislike <- lm(log_dislike ~ funny + show_product_quickly + patriotic +
               celebrity + danger + animals + use_sex + year,
               data = yt_new)
summary(m_dislike)
```

Call:

```
lm(formula = log_dislike ~ funny + show_product_quickly + patriotic +
    celebrity + danger + animals + use_sex + year, data = yt_new)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-4.0292	-1.3299	-0.3192	0.8986	8.7219

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-183.06813	53.34768	-3.432	0.000719 ***
funnyTRUE	0.25949	0.35154	0.738	0.461224
show_product_quicklyTRUE	0.27511	0.30593	0.899	0.369515
patrioticTRUE	0.81407	0.41859	1.945	0.053095 .
celebrityTRUE	-0.20214	0.32125	-0.629	0.529852
dangerTRUE	0.22184	0.31400	0.707	0.480630
animalsTRUE	-0.21192	0.29675	-0.714	0.475911
use_sexTRUE	-0.32980	0.33681	-0.979	0.328583
year	0.09207	0.02653	3.471	0.000626 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.074 on 216 degrees of freedom

(22 observations deleted due to missingness)

Multiple R-squared: 0.09753, Adjusted R-squared: 0.06411

F-statistic: 2.918 on 8 and 216 DF, p-value: 0.004115

```
# 4. Model for Comment Counts
```

```
m_comment <- lm(log_comment ~ funny + show_product_quickly + patriotic +
  celebrity + danger + animals + use_sex + year,
  data = yt_new)
summary(m_comment)
```

Call:

```
lm(formula = log_comment ~ funny + show_product_quickly + patriotic +
  celebrity + danger + animals + use_sex + year, data = yt_new)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-4.1372	-1.4665	-0.1427	1.4061	5.8468

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-99.09835	52.92351	-1.872	0.0625 .
funnyTRUE	0.21954	0.34528	0.636	0.5256
show_product_quicklyTRUE	0.40939	0.30229	1.354	0.1771
patrioticTRUE	0.66698	0.39902	1.672	0.0961 .
celebrityTRUE	0.29767	0.31541	0.944	0.3464
dangerTRUE	0.17784	0.31069	0.572	0.5677
animalsTRUE	-0.26802	0.29347	-0.913	0.3621
use_sexTRUE	-0.39323	0.33163	-1.186	0.2370
year	0.05034	0.02632	1.913	0.0571 .

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.039 on 213 degrees of freedom

(25 observations deleted due to missingness)

Multiple R-squared: 0.06535, Adjusted R-squared: 0.03025

F-statistic: 1.862 on 8 and 213 DF, p-value: 0.06748

I ran four linear regression models using the log of views, likes, dislikes, and comments as the outcomes. The predictors were the seven ad features and the year of the ad. Most ad features were not statistically significant in any of the models. The main consistent result was that the variable year was positive and significant for likes and dislikes, and nearly significant for comments. This means that YouTube engagement increased over time, which makes sense because the platform grew in popularity. For views, no variables were significant.

### Problem 3d

```
yt_new$log_view <- log1p(yt_new$view_count)

#Selecting the needed columns (outcome + predictors)
view_data <- yt_new[, c("log_view", "funny", "show_product_quickly",
                        "patriotic", "celebrity", "danger",
                        "animals", "use_sex", "year")]

#Dropping rows with missing values
view_data <- na.omit(view_data)

#Converting true/false flags to 0 and 1
view_data$funny <- as.integer(view_data$funny)
view_data$show_product_quickly <- as.integer(view_data$show_product_quickly)
view_data$patriotic <- as.integer(view_data$patriotic)
view_data$celebrity <- as.integer(view_data$celebrity)
view_data$danger <- as.integer(view_data$danger)
view_data$animals <- as.integer(view_data$animals)
view_data$use_sex <- as.integer(view_data$use_sex)

#Creating y (outcome) and X (design matrix with intercept)
y <- as.matrix(view_data$log_view)
X <- model.matrix(~ funny + show_product_quickly + patriotic +
                  celebrity + danger + animals + use_sex + year,
                  data = view_data)

#Applying OLS formula: beta = (X'X)^(-1) X'y
XtX <- t(X) %*% X
Xty <- t(X) %*% y
beta_hat <- solve(XtX) %*% Xty # Or use solve(XtX, Xty)
cat("Manual beta calculation:\n")
```

Manual beta calculation:

```
print(beta_hat)
```

```
              [,1]
(Intercept) -31.55015804
funny        0.56492445
show_product_quickly 0.21088918
patriotic    0.50699051
celebrity    0.03547862
danger       0.63131085
animals     -0.31001838
use_sex      -0.38670726
year         0.02053399
```

```
#Comparing with lm()
m_view <- lm(log_view ~ funny + show_product_quickly + patriotic +
             celebrity + danger + animals + use_sex + year,
             data = view_data)

cat("\nLM function coefficients:\n")
```

LM function coefficients:

```
print(coef(m_view))
```

```
      (Intercept)      funny show_product_quickly
-31.55015804      0.56492445      0.21088918
      patriotic      celebrity      danger
0.50699051      0.03547862      0.63131085
      animals      use_sex      year
-0.31001838     -0.38670726      0.02053399
```

```
#Verifying they are identical
cat("\nDifference between manual and lm:\n")
```

Difference between manual and lm:

```
difference <- beta_hat - coef(m_view)
print(difference)
```

```

                                [,1]
(Intercept) -4.830980e-11
funny        -2.831069e-13
show_product_quickly 3.930190e-14
patriotic     8.892886e-14
celebrity     7.675249e-13
danger       -1.365574e-14
animals      1.273981e-13
use_sex      -1.317835e-13
year         -1.374387e-13
```

```
cat("\nMaximum difference:", max(abs(difference)), "\n")
```

```
Maximum difference: 4.83098e-11
```

I calculated the regression coefficients for view counts manually. The manual results matched the coefficients from the `lm()` function. The maximum difference between the two sets of coefficients was only  $4.83 \times 10^{-11}$  which is almost zero. This confirms that the manual matrix algebra approach and the `lm()` function give the same result.

## Attribution of Sources

For Problem 1, I used the following references:

<https://www.rdocumentation.org/packages/stats/versions/3.6.2/topics/Uniform> <https://stat.ethz.ch/R-manual/R-devel/library/stats/html/Uniform.html> - I used these to learn how to generate random numbers with `runif()` which I needed to decide whether each step should be +10 or -3. <https://r4ds.hadley.nz/functions.html> - I used this to understand how to write my own functions in R. <https://stackoverflow.com/questions/21991130/simulating-a-random-walk> - I used this for ideas on how to set up a random walk simulation in R.

For Problem 2, I used the following reference: <https://stat.ethz.ch/R-manual/R-devel/library/stats/html/Poisson.html> <https://stat.ethz.ch/R-manual/R-devel/library/stats/html/Normal.html> <https://bstaton1.github.io/au-r-workshop/ch4.html> - I used this to understand Monte Carlo simulation <https://stat.ethz.ch/R-manual/R-devel/library/base/html/matrix.html> - I referred to this to understand matrices

For Problem 3, I used the following references: <https://r4ds.had.co.nz/vectors.html#subsetting-1>  
<https://stat.ethz.ch/R-manual/R-devel/library/stats/html/model.matrix.html> - I used this to create the design matrix <https://r-statistics.co/Linear-Regression.html> - I used this to understand how linear regression works

## **Github Repository**

<https://github.com/prathii7/Computational-Methods-and-Tools-in-Statistics>